

Cloud Computing Technology
CSC / ECE - 547 Fall 2023

Global Expansion of Fifth Third Bank

Kartik Soni - ksoni, Vansh Mehta - vpmehta2

Plagiarism Declaration

We, the team members, understand that copying & pasting material from any source in our project is an allowed practice; we understand that not properly quoting the source constitutes plagiarism.

All team members attest that we have properly quoted the sources in every sentence/paragraph we have copy & pasted in our report. We further attest that we did not change words to make copy & pasted material appear as our work.

Table of Contents	Page No.
1. Introduction	
1.1 Motivation	4
1.2 Executive Summary	4
2. Problem Description	
2.1 The Problem	4
2.2 Business Requirements	4
2.3 Technical Requirements	5
2.4 Tradeoffs and Conflicting Requirements	7
3. Provider Selection	
3.1 Criteria for choosing a provider	7
3.2 Provider Comparison	9
3.3 The final selection	12
3.3.1 List of services	13
3.4 The list of services offered by the winner	13
4. The first design draft	
4.1 Basic building blocks of the design	14
4.2 Top-level informal validation	16
4.3 Action items and rough timeline	16
5. Second Design draft	
5.1 Use of Well-Architected Framework	17
5.2 Discussion of Pillars	19
5.3 Use of CloudFormation diagrams	20
5.4 Validation of the design	20
5.5 Design principles and best practices used	21
5.6 Tradeoffs revisited	23
5.7 Discussion of an alternate design	25
6. Kubernetes Experimentation	
6.1 Experimental Design	25
6.2 Workload Generation with Locust	28
6.3 Analysis of the results	31
7. Ansible Playbooks	
8. Demonstration	
9. Comparisons	

10. Conclusions	
10.1 Lessons Learned	34
10.2 Possible continuation of the project	34
11. References	34

1. Introduction

1.1 Motivation

Our motivation for this project is to enable the global expansion of a small bank, specifically Fifth Third Bank. As cloud architects, we aim to harness the power of cloud computing to facilitate the bank's expansion into international markets. By doing so, we can help the bank unlock new opportunities, serve a broader customer base, and compete effectively in the global financial landscape.

1.2 Executive Summary

In this project, our goal is to architect a cloud-based solution to support the global expansion of Fifth Third Bank. This summary is intended for stakeholders involved in the bank's expansion efforts. By leveraging cloud technology, we aim to ensure global accessibility, high availability, security and compliance, and seamless scalability for the bank's services. We will address issues like data sovereignty, disaster recovery, cross-border transactions, multi-currency support, and language localization to meet the diverse needs of international customers. Our solution will also focus on regulatory compliance, cost optimization, and environmentally sustainable practices. The cloud architect's role is to design an infrastructure that aligns with the business requirements for a successful global expansion of Fifth Third Bank.

2. Problem Description

2.1 The problem

The high-level problem we are addressing is the global expansion of Fifth Third Bank. The bank, originally serving a local or regional market, is now seeking to operate on a global scale. This expansion presents numerous challenges and opportunities, including compliance with diverse international regulations, the need for 24/7 customer support across different time zones, and the demand for seamless cross-border transactions. The cloud architect's role is to design an IT infrastructure that can support and enable this global expansion while ensuring high availability, data security, and efficient cost management.

2.2 Business Requirements

BR1: Ensure that the application can be accessed from anywhere in the world with low latency.

BR2: Ensure the application complies with country-specific regulations such as data residency, cross-border data transfers, data retention, and others.

BR3: High Availability: Guarantee uninterrupted access to banking services.

BR4: Scalability: Ensure the platform can seamlessly handle increased customer demand during peak hours or new market entries.

BR5: Security and Compliance: Implement robust security measures to protect customer data and meet international compliance standards like GDPR [31], CCPA [32], and local financial regulations.

BR6: Ensure disaster recovery strategies to safeguard customer data in the event of unforeseen incidents.

BR7: Load Balancing: Distribute incoming traffic efficiently to maintain performance during high demand.

BR8: Global Data Redundancy: Replicate and backup customer data across multiple geographic regions.

BR9. Cost Optimization: Implement cost control measures to manage cloud infrastructure expenses effectively.

BR10. Use cloud-based auto-scaling to allocate resources dynamically as needed.

BR11. Ensure unauthorized access is detected and caught.

BR12. Ensure a virtual private network for employees for very low latency and protection against network outages.

BR13: Real-time Monitoring and Alerting: Establish a real-time monitoring and alerting system to proactively detect and respond to issues related to cloud resources.

BR14: Identify Tenants in our cloud space as the users.

2.3 Technical Requirements

[BR1] TR1.1 & 7.1: Geographic Load Balancer: Configure load balancing to distribute incoming traffic efficiently across multiple geographic regions to minimize latency and ensure global access.

TR2.1: Develop a plan for managing data residency requirements by storing data in data centers within specific regions or countries to comply with local regulations.

TR2.2: Develop a strategy for cross-border data transfers to meet local data compliance requirements.

TR2.3: Create data retention policies to comply with local data retention regulations.

TR3.1: Multi-Region Failover Architecture - Design a multi-region architecture with automatic failover capabilities to ensure uninterrupted access to banking services in case of regional outages or disasters.

TR4.1: Auto-Scaling Policies - Define policies for dynamic resource allocation based on metrics like CPU utilization or incoming traffic to automatically adjust resources to meet performance demands.

TR5.1: Security Group Design - Design security groups for virtual networks to control inbound and outbound traffic, addressing both security and compliance requirements.

[BR6] TR6.1 & 8.1: Regional Redundancy Strategy - Deploy redundant resources (including data and elastic compute services) across multiple regions to ensure data redundancy and minimize downtime.

TR6.2: Automate backup and data recovery strategies.

[BR7] TR1.1 & 7.1: Geographic Load Balancer: Configure load balancing to distribute incoming traffic efficiently across multiple geographic regions to minimize latency and ensure global access.

[BR8] TR 6.1 & 8.1: Regional Redundancy Strategy - Deploy redundant resources (including data and elastic compute services) across multiple regions to ensure data redundancy and minimize downtime.

TR9.1: Establish a consistent resource tagging policy for tracking costs and resource ownership across the cloud environment.

TR10.1: Configure auto-scaling policies based on metrics like CPU utilization or incoming traffic to automatically adjust resource allocation.

TR11.1: Design the security group structure with identity and access management services for virtual networks to control inbound and outbound traffic, addressing unauthorized access attempts.

TR12.1: Establish a Virtual Private Network (VPN) solution to provide low-latency, secure access for employees.

TR13.1: Configure real-time monitoring and alerting to detect and respond to issues related to cloud resources.

TR14.1: Accurately identifies and segregates tenant data, ensuring that each tenant's data is isolated and managed according to the specific regulatory requirements of their region.

TR14.2: Utilizes identity and access management (IAM) [7] solutions to assign and manage credentials and access policies for each tenant, guaranteeing that only authorized users can access their respective data and services.

2.4 Tradeoffs

Trade-off Set 1: BR4 vs. (BR9 and BR10)

Reason: BR4 (Scalability) aims to ensure the platform can seamlessly handle increased customer demand during peak hours or new market entries, which may require additional resources.

On the other hand, BR9 and BR10 together seek to optimize costs by adjusting resources dynamically. These two objectives can sometimes conflict because scaling for peak demand may result in higher costs.

Trade-off Set 2: BR5 vs. (BR6 and BR9)

The trade off involves balancing the redundancy strategy's benefits in terms of reliability and minimizing downtime with the costs associated with maintaining and synchronizing data across multiple regions. Meeting regional compliance requirements may also influence the choice of regions for data redundancy. Finding the right balance between these factors is essential.

3. Provider Selection

Selecting the right cloud provider is a critical decision that will have long-term implications on the bank's ability to scale, innovate, and maintain a secure and compliant global operation. The following criteria will serve as our guide to evaluate the three major cloud service providers: Amazon Web Services (AWS), Microsoft Azure [33], and Google Cloud Platform (GCP) [34].

3.1 Criteria for choosing a provider [1]

1. **Comprehensive Service Offerings:** The provider must have a broad range of services that align with our Business and Technical Requirements (BRs and TRs), such as geographic load balancing, auto-scaling, compliance, and security services.
2. **Global Reach and Data Center Locations:** Providers must have a vast global network of data centers that offer low-latency access and data residency options to comply with various international laws.
3. **Compliance and Security:** The ability to meet international compliance standards like GDPR, CCPA, and financial industry regulations is essential. Providers must also offer robust security measures that align with our security and compliance requirements.

4. High Availability and Reliability: Providers should guarantee high uptime percentages and offer multi-region failover capabilities to ensure uninterrupted service.
5. Scalability and Performance: The cloud provider must offer solutions that can scale resources up or down dynamically in response to our operational demands and performance metrics.
6. Cost Management and Optimization: We need tools that help monitor, manage, and optimize costs related to cloud resource usage.
7. Data Management Capabilities: Capabilities for managing data residency, retention, and cross-border data transfers are vital to meet BR2 and associated TRs.
8. Disaster Recovery and Redundancy: Providers should offer comprehensive disaster recovery solutions and data redundancy across regions to maintain service continuity.
9. Security and Identity Management: Advanced security and identity management solutions are required to ensure that only authorized users have access to specific cloud resources.
10. Network Infrastructure and Performance: The cloud provider must have robust networking capabilities, including a VPN solution for secure employee access, low latency, and protection against network outages.
11. Operational Management: Tools for real-time monitoring, alerting, and reporting are crucial for proactive issue resolution and maintaining operational health.
12. Customer Support and SLAs: The level of customer support, including the availability of cloud support engineers and the details of the Service Level Agreements (SLAs), will be crucial, especially for critical financial operations.
13. Innovation and Ecosystem: Continuous improvement in service offerings and a rich ecosystem of partners and services are important for future-proofing our infrastructure.
14. Experience in the Financial Sector: The provider's track record and expertise in handling financial sector clients will be taken into consideration.
15. Ease of Migration and Integration: The tools and services that support a smooth migration from our current setup to the cloud and integration with our existing systems are necessary.

3.2 Provider comparison

The following is a table providing a brief comparison [1]

Criteria	AWS	Azure	GCP	Justification
Comprehensive Service Offerings	1	2	3	AWS offers the widest range of services, Azure follows, and GCP is third, though all are rapidly expanding their offerings.
Global Reach	1	2	3	AWS has the largest global presence, Azure is close behind, while GCP is growing but has fewer regions compared to the others.
Compliance and Security	2	1	3	Azure often leads in enterprise compliance solutions, AWS is close in offering, while GCP is slightly behind.
High Availability and Reliability	1	2	3	AWS has the longest track record, with Azure and GCP following but improving steadily.

Scalability and Performance	1	2	3	AWS's maturity gives it the edge in scalability and performance, though Azure and GCP are competitive.
Cost Management and Optimization	2	3	1	GCP often offers more aggressive discounts and flexible contracts, AWS and Azure are more traditional.
Data Management Capabilities	1	2	3	AWS provides more mature tools for data management, Azure is close, and GCP is in third place.
Disaster Recovery and Redundancy	1	2	3	AWS's maturity again provides a slight edge, with Azure and GCP not far behind.
Security and Identity Management	2	1	3	Azure integrates well with existing Microsoft-based security tools, which many enterprises use.
Network Infrastructure and Performance	2	1	3	Azure's integration with Microsoft's enterprise tools gives it an edge for some companies, followed by AWS and then GCP.

Operational Management	1	2	3	AWS's suite of management tools is quite extensive, with Azure and GCP following.
Customer Support and SLAs	1	2	3	AWS has a long-standing reputation for strong support and comprehensive SLAs, with Azure and GCP in pursuit.
Innovation and Ecosystem	1	3	2	AWS is known for rapid innovation, GCP is recognized for AI and analytics, and Azure is strong but more conservative.
Experience in the Financial Sector	2	1	3	Azure benefits from Microsoft's longstanding enterprise and financial sector relationships, AWS and GCP follow.
Ease of Migration and Integration	2	1	3	Azure provides seamless integration for existing Microsoft customers, while AWS has more general services, and GCP is third.

Justification:

- AWS is the market leader with the most extensive portfolio of services.
- AWS has the most extensive global network, with Azure growing rapidly and GCP expanding strategically.
- Azure is often chosen by enterprises for its strong compliance offerings, particularly for companies with existing Microsoft software compliance dependencies.
- AWS has set the standard with its established infrastructure.

- AWS's maturity and breadth of options give it an advantage.
- GCP is known for competitive pricing and customer-friendly cost management tools.
- AWS has a strong track record with its wide array of database services.
- AWS's robustness and proven disaster recovery capabilities give it an edge.
- AWS's integration with Amazon's security and identity services is a significant advantage.
- Azure's core networking services are highly regarded, especially for organizations that are already using Microsoft's enterprise software.
- AWS's operations management tools and experience are well-regarded in the industry.
- AWS has historically set a high standard for customer support.
- AWS is recognized for its pace of innovation, GCP for its strengths in AI and machine learning, and Azure for its integration with Microsoft's broad suite of tools.
- Azure and AWS both have strong financial sector clientele, but Azure's deep integration with other Microsoft products gives it a lead in this domain.
- Azure's compatibility with other Microsoft products can make it the path of least resistance for enterprises already invested in the Microsoft ecosystem.

3.3 The final selection

After a thorough analysis and comparison of the top cloud service providers, we have concluded that **AWS** is the best fit for Fifth Third Bank's global expansion needs as outlined in the technical requirements (TRs).

AWS emerges as the clear winner for several reasons. Firstly, its comprehensive service offerings are unmatched, which aligns with our need for a diverse range of services to meet our specific TRs. AWS's breadth of services ensures that we can leverage the latest technologies and innovations without the need to engage with multiple providers.

Secondly, AWS's global reach with the largest number of data centers around the world supports our requirement for low-latency access (BR1) and aligns with our need for global data redundancy (BR8). This expansive network is essential to ensure high availability (BR3) and disaster recovery capabilities (BR6).

The high level of security and compliance standards maintained by AWS aligns with our security and compliance requirement (BR5), offering robust measures to protect customer data and comply with international standards like GDPR [31] and CCPA [32].

In terms of scalability (BR4), AWS's auto-scaling services and elastic load balancing capabilities (TR4.1 & TR7.1) are industry-leading, ensuring that our platform can scale seamlessly with changing customer demands.

Cost optimization (BR9) is also a critical factor, and AWS provides a comprehensive set of tools to control and optimize costs. The ability to use auto-scaling (BR10) effectively allows us to

manage expenses while maintaining performance, a delicate balance that AWS has been proven to handle effectively.

Operational management and customer support are other areas where AWS excels. Its suite of management tools, combined with its customer support and service level agreements (SLAs), provide a safety net that is crucial for our operations.

While Azure and GCP both offer compelling features, particularly in areas of compliance and security for Azure and data analytics and machine learning for GCP, AWS's established leadership and proven track record in supporting financial institutions make it the superior choice for our bank's requirements. Therefore, we select AWS as our cloud service provider to enable Fifth Third Bank's global expansion, confident that it provides the best combination of services, performance, security, and cost efficiency to meet our ambitious project goals.

3.3.1 List of services

1. Amazon EC2 (Elastic Compute Cloud) - Virtual servers for scalable computing capacity. [3]
2. Amazon S3 (Simple Storage Service) - Scalable object storage for data backup, archival and analytics. [4]
3. Amazon RDS (Relational Database Service) - Managed relational database service for MySQL, PostgreSQL, Oracle, SQL Server, and MariaDB. [5]
4. Amazon VPC (Virtual Private Cloud) - A private section of the AWS cloud for launching AWS resources in a virtual network. [6]
5. AWS IAM (Identity and Access Management) - Secure access management for AWS services and resources. [7]
6. Amazon CloudFront - A content delivery network (CDN) service for delivering data, videos, applications, and APIs. [8]
7. AWS Lambda - Event-driven, serverless computing service. [9]
8. AWS Auto Scaling - Monitoring and automatic scaling of resources. [10]
9. AWS CloudTrail - Service that provides logs of all user activity and API usage. [11]
10. Amazon Route 53 - A scalable and highly available Domain Name System (DNS) web service. [12]
11. AWS Shield - Managed Distributed Denial of Service (DDoS) protection. [13]
12. Amazon GuardDuty - A threat detection service that continuously monitors for malicious activity. [14]
13. AWS Key Management Service (KMS) - Managed service that makes it easy to create and control encryption keys used to encrypt data. [15]
14. Amazon CloudWatch - Monitoring service for AWS cloud resources and applications. [16]
15. Amazon ElastiCache - For in-memory caching to decrease latency and improve the speed of applications. [17]

4 The First Design Draft

In constructing the initial design draft for Fifth Third Bank's global cloud architecture, several fundamental components and characteristics stand out. This draft will highlight how these components interlock to create a robust, scalable, and secure system.

4.1 The Basic Building Blocks of the Design

The proposed cloud infrastructure is based on the following foundational elements:

- Multi-Region Deployment: Utilize AWS's global infrastructure to deploy services across multiple regions, ensuring high availability and disaster recovery.
- Compute:
 - Utilize AWS EC2 instances configured with Auto Scaling groups to dynamically adjust capacity to maintain steady, predictable performance at the lowest possible cost.
 - Implement Elastic Load Balancing (ELB) to automatically distribute incoming application traffic across multiple targets, such as EC2 instances, in multiple Availability Zones, which increases the fault tolerance of your applications.
 - Integrate AWS Lambda for serverless computing, allowing for the execution of code in response to events and automatically managing the underlying compute resources.
- Storage: Leverage Amazon S3 for scalable object storage, with multi-region redundancy and lifecycle policies aligned with data residency requirements.
- Database: Employ Amazon RDS and Amazon DynamoDB for relational and NoSQL database needs, respectively, ensuring high performance and automated backups.
- Networking: Establish a Virtual Private Cloud (VPC) for secure and isolated network configuration, with AWS Transit Gateway to simplify inter-region connectivity.
- Security: Implement AWS Identity and Access Management (IAM) for fine-grained access control and utilize AWS Shield for DDoS mitigation.
- Compliance: Enforce data compliance using AWS Config and AWS Key Management Service (KMS) to manage encryption keys.
- Monitoring and Management: Adopt Amazon CloudWatch and AWS CloudTrail for monitoring, logging, and real-time alerting of operational and security events.
- Cost Management: Use AWS Cost Explorer and AWS Budgets to monitor and optimize expenses.

Phase 1: Core Infrastructure Setup

VPC Configuration: Establish a Virtual Private Cloud (VPC) with public and private subnets across multiple Availability Zones (AZs) to ensure high availability and fault tolerance.

Basic EC2 Setup: Launch EC2 instances in each AZ without auto-scaling, just to ensure that the most basic form of compute resource is available.

Phase 2: High Availability and Scalability

Elastic Load Balancing (ELB): Introduce an Elastic Load Balancer to distribute incoming traffic evenly across the EC2 instances in different AZs.

Auto Scaling Groups: Implement Auto Scaling Groups (ASGs) for EC2 instances within each AZ to automatically adjust the number of EC2 instances based on demand.

Phase 3: Service Orchestration

EKS Introduction: Replace standalone EC2 instances with an Elastic Kubernetes Service (EKS) cluster for better orchestration of services.

Service Deployment: Deploy essential services within EKS, initially perhaps without ingress controllers or sophisticated routing.

Phase 4: Security and Compliance

IAM Policies: Define Identity and Access Management (IAM) policies to enforce security best practices and grant appropriate permissions.

Monitoring and Auditing: Set up CloudTrail for auditing API calls and GuardDuty for threat detection.

Phase 5: Database and Storage

RDS Setup: Provision a managed Relational Database Service (RDS) for structured data storage.

S3 Integration: Integrate Amazon S3 for scalable object storage, ensuring that data is available across services.

Phase 6: Microservices and Serverless

Lambda Functions: Incorporate AWS Lambda for serverless computations, reducing the need for constant EC2 compute.

Service Decoupling: Begin decoupling services, potentially using Amazon SNS for pub/sub messaging to coordinate microservices.

Phase 7: Performance Optimization

Elasticache: Add Elasticache to the mix to enhance performance through in-memory caching of frequently accessed data.

Advanced Monitoring: Utilize CloudWatch along with custom metrics and alarms to closely monitor the performance and health of the system.

Phase 8: Networking and DNS

Route 53 Configuration: Implement Route 53 for DNS management, connecting user requests to the ELB efficiently.

Advanced Networking: Refine NAT Gateway setup and routing rules for traffic management between public and private subnets.

Phase 9: Cost Management and Optimization

Cost Analysis Tools: Deploy Cost Explorer and Budgets to monitor and manage AWS costs, ensuring the architecture remains within financial constraints.

Phase 10: Security Enhancement

AWS Shield: Employ AWS Shield Standard for protection against DDoS attacks, enhancing the overall security posture.

Phase 11: Final Touches and Integration

Ingress Control: Finalize the architecture by setting up Ingress control within EKS for efficient routing to backend services.

Final Testing: Conduct extensive load testing and security assessments to ensure the architecture meets all functional and non-functional requirements.

4.2 Top-Level, Informal Validation of the Design [1]

This initial design blueprint offers a high-level assurance that the TRs will be met:

- Global Accessibility and Low Latency: The use of multi-region deployment and Amazon CloudFront for content delivery is expected to meet global accessibility needs and reduce latency.
- Data Compliance and Residency: Amazon S3's cross-region replication capabilities combined with local data residency options in RDS and DynamoDB aim to comply with country-specific regulations.
- High Availability and Disaster Recovery: The multi-region deployment and automated failover strategies inherent in AWS services are designed to offer robust availability and disaster recovery solutions.
- Scalability: The auto-scaling features of EC2 and Lambda are expected to provide the necessary scalability to manage fluctuating demands and grow with the bank's expansion.
- Security and Compliance: AWS's comprehensive security and compliance services, including IAM, Shield, KMS, and Config, provide a strong argument for the design's capability to protect sensitive data and adhere to regulatory standards.

4.3 Action Items and Rough Timeline

Skipped

5 The Second Design Draft

5.1 Use of the Well-Architected Framework

To refine our design for the global cloud architecture of Fifth Third Bank, we adhere to the AWS Well-Architected Framework's structured approach. This approach offers guidance to build secure, high-performing, resilient, and efficient infrastructure for applications. Below are the distinct steps suggested by the framework: [2]

1. **Review and Benchmarking:** Begin with an assessment of the current architecture against the best practices of the Well-Architected Framework. This step involves using the AWS Well-Architected Tool that provides a consistent process for evaluating architectures and implementing designs that will scale over time.
2. **Define Workload:** Identify and define the workload within the AWS environment. A workload is any application or set of applications and their associated configuration, deployed as code on AWS infrastructure. Understanding the workload's specific needs is crucial to aligning it with the correct architectural choices.
3. **Design Principles:** Apply the design principles associated with each pillar of the Well-Architected Framework. These principles guide the design and deployment of the workload in a scalable and adaptable manner. They include strategies like implementing disposable resources instead of fixed servers, automating to make architectural experimentation easier, and building with the assumption that everything will fail at some point.
4. **Five Pillars:** Evaluate the architecture against the five original pillars of the framework: Operational Excellence, Security, Reliability, Performance Efficiency, and Cost Optimization. Each pillar has a set of best practices and key concepts to apply.
5. **Sustainability:** Integrate the newly added sixth pillar, Sustainability, into the architecture. This involves assessing how to improve energy efficiency and reduce the carbon footprint of the workload by leveraging AWS's infrastructure.
6. **Implement Strategies:** Based on the review and the framework's guidance, implement strategies that address any identified deficiencies. This could include re-architecting for greater scalability, adding security controls, automating manual processes, or optimizing resources to reduce costs.
7. **Measure and Improve:** Continuously measure the architecture's performance against the established benchmarks and improve upon it. The AWS Well-Architected Framework is iterative, and as the workload or business requirements change, the architecture should evolve accordingly.

By following the structured process suggested by the AWS Well-Architected Framework, we aim to create a cloud infrastructure that is not only aligned with current best practices but also positioned for future growth and sustainability.

The WAF is a strategic set of best practices and guidelines provided by AWS tailored to support the global expansion project of Fifth Third Bank. This comprehensive framework assists in building and maintaining a secure, high-performing, resilient, and efficient cloud architecture. By adhering to the principles outlined in the framework, the bank can systematically evaluate its workloads, identify improvement areas, and ensure that its cloud architecture aligns with

industry best practices. The framework is organized into six pillars, each addressing crucial aspects relevant to the global expansion initiative: [1]

Operational Excellence:

Objective: Effective management of development and operational workloads with continuous improvement for enhanced business value through efficient processes and procedures.

Related TRs:

- a. TR1.1 (Geographic Load Balancer): Implementing robust user authentication and authorization to ensure secure operational practices.
- b. TR1.2 (Unique IDs for Users): Assigning unique IDs to every user streamlines user management and enhances security.

Security:

Objective: Addressing data protection, system security, and overall asset security to fortify security postures and mitigate potential risks.

Related TRs:

- a. TR8.1 (Security Group Design): Implementing encryption mechanisms for secure storage of uploaded documents prioritizes data security.
- b. TR8.2 (Security Group Design): Utilizing access control prevents unauthorized access, reinforcing overall document security.

Reliability:

Objective: Ensuring workloads consistently perform their intended functions correctly, providing best practices for operating and testing workloads throughout their lifecycle.

Related TRs:

- a. TR6.1 (Auto-Scaling Policies): Automated backups of the database prevent data loss, ensuring reliability in data storage.
- b. TR6.2 (Auto-Scaling Policies): Establishing a disaster recovery plan with backup and restoration procedures enhances overall workload reliability.

Performance Efficiency:

Objective: Adept use of computing resources to meet system requirements efficiently, adapting to changes in demand while maintaining optimal performance.

Related TRs:

- a. TR2.1 (Data Residency Plan): Comprehensive monitoring for real-time application health optimizes performance efficiency.
- b. TR5.1 (Auto-Scaling Policies): Auto-scaling configurations dynamically adjust resources, ensuring efficient handling of document upload demands.

Cost Optimization:

Objective: Focused on delivering business value at the most economical price point, emphasizing the importance of managing costs effectively.

Related TRs:

- a. TR5.1 (Auto-Scaling Policies): Auto-scaling configurations optimize resource usage based on document upload demands, contributing to cost efficiency.
- b. TR7.2 (Load Balancing Mechanisms): Load balancing mechanisms distribute traffic for cost-efficient resource utilization.

Sustainability:

Objective: Highlighting the ongoing commitment to reduce energy consumption, enhance efficiency, and minimize environmental impact across all aspects of the workload.

Related TRs:

- a. TR5.1 (Auto-Scaling Configurations): Leverage auto-scaling configurations and load balancing mechanisms to optimize resource utilization, indirectly promoting energy efficiency.

5.2 Discussion of pillars

For the global cloud architecture of Fifth Third Bank, two pillars from the AWS Well-Architected Framework are particularly pivotal: Security and Reliability. These pillars ensure that the bank's infrastructure is not only protected against various threats but is also dependable for continuous financial operations. [2]

Security: The Security pillar is of utmost importance for any financial institution, given the sensitive nature of financial data and the stringent regulatory requirements. The best practices under this pillar ensure that data, systems, and assets are well-protected. For our project, implementing a strong identity foundation is crucial. We will enforce the principle of least privilege and embed security in every aspect of the architecture. Encryption, both at rest and in transit, will be a default approach, safeguarding data integrity and confidentiality.

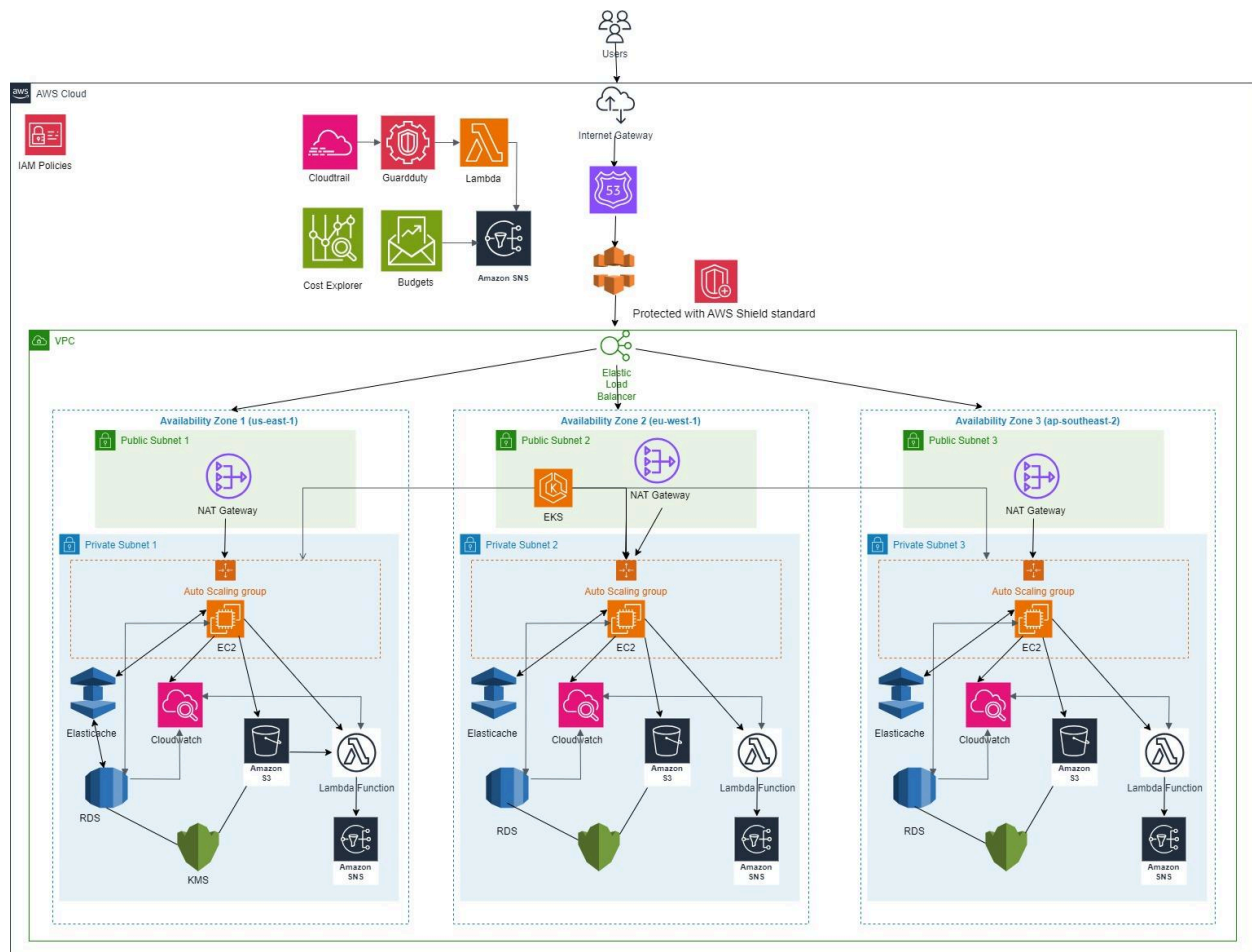
Key AWS services that support the Security pillar include AWS Identity and Access Management (IAM) for managing access, Amazon VPC for secure networking, AWS Key Management Service (KMS) for managing cryptographic keys, and AWS Shield for DDoS protection. Implementing these services will help create a comprehensive defense strategy that addresses various aspects of security, from data encryption to network isolation and access control.

Reliability: The Reliability pillar is critical as it ensures the banking system's ability to perform its intended functions correctly and consistently. For Fifth Third Bank, this means ensuring that the financial transactions and operations are always available to customers and staff. The architecture must automatically recover from failures, dynamically acquire computing resources to meet demand, and mitigate disruptions such as network issues or DDoS attacks.

To support the Reliability pillar, we will leverage AWS services like Amazon EC2 Auto Scaling to manage the compute capacity and maintain application availability. Amazon RDS and Amazon DynamoDB with cross-region replication will provide data redundancy. AWS CloudFormation will be used to manage and provision resources in an orderly and predictable fashion. AWS Elastic

Load Balancing will distribute incoming application traffic across multiple targets, such as EC2 instances, in different Availability Zones, which increases the fault tolerance of the application.

5.3 Use of Cloudformation Diagrams



5.4 Validation of the Design

This section will discuss how the template meets these requirements.

Meeting Business Requirements (BRs) and Technical Requirements (TRs):

1. **Global Accessibility and Low Latency (BR1, TR1.1 & 7.1):** The architecture leverages Amazon EC2 with Auto Scaling and Elastic Load Balancing (ELB) across multiple Availability Zones. This setup ensures that the application can handle incoming traffic efficiently, thus minimizing latency and ensuring global accessibility.

2. **Data Compliance and Residency (BR2, TR2.1 - TR2.3):** By deploying Amazon RDS, the architecture allows for data storage in region-specific data centers, addressing local data residency requirements. This setup is crucial for complying with regulations like GDPR.
3. **High Availability (BR3, TR3.1):** The multi-region deployment strategy, coupled with Amazon S3's cross-region replication and RDS's multi-AZ deployment, ensures high availability and robust disaster recovery capabilities.
4. **Scalability (BR4, TR4.1):** Auto-scaling policies defined in EC2 instances and the use of AWS Lambda provide dynamic scaling capabilities to efficiently handle varying loads.
5. **Security and Compliance (BR5, TR5.1):** Security Groups and IAM roles integrated within the architecture ensure fine-grained access control, enhancing the overall security posture.
6. **Disaster Recovery (BR6, TR6.1 & 8.1):** The replication of S3 buckets and multi-region deployment of key resources ensure data redundancy and recovery capabilities.
7. **Cost Optimization (BR9, TR9.1):** The use of AWS Cost Explorer and AWS Budgets, alongside efficient resource utilization strategies like auto-scaling, enables effective cost management.
8. **Auto-Scaling (BR10, TR10.1):** Auto-scaling groups in EC2 instances are directly responsive to changes in demand, ensuring resource availability while optimizing costs.
9. **Unauthorized Access Detection (BR11, TR11.1):** Security groups, along with AWS IAM, provide robust mechanisms to detect and prevent unauthorized access attempts.
10. **VPN for Employees (BR12, TR12.1):** While not explicitly defined in the template, the architecture's design allows for the integration of VPN solutions to ensure secure and low-latency connections for employees.
11. **Real-time Monitoring (BR13, TR13.1):** The implementation of Amazon CloudWatch and AWS CloudTrail ensures real-time monitoring and alerting, enabling proactive issue resolution.
12. **Tenant Identification (BR14, TR14.1 & 14.2):** The architecture's multi-tenant design, supported by AWS's IAM and database services, ensures effective data segregation and access control for different tenants.

5.5 Design Principles and Best Practices Used

The design of the Fifth Third Bank cloud infrastructure is based on a combination of AWS Well-Architected Framework (WAF) principles and cloud architecture best practices. Here we outline how these principles are applied within the context of the bank's business and technical requirements (BRs and TRs).

Principle Application in Design:

Operational Excellence: By automating deployments with AWS CloudFormation, we ensure consistent and repeatable processes, a key aspect of operational excellence. This is reflected in

TR3.1, TR4.1, and TR10.1, where automation in failover, scaling, and resource allocation are critical.

Security: The principle of security by design is evident in the architecture through the use of AWS Identity and Access Management (IAM) and security groups. This directly addresses TR5.1 and BR5, ensuring fine-grained access controls and protection against unauthorized access.

Reliability: We applied the principle of reliability through a multi-region deployment and Auto Scaling. This addresses BR3 and TR3.1, ensuring the system remains operational and maintains its performance despite potential disruptions.

Performance Efficiency: The use of Auto Scaling groups for EC2 instances and Elastic Load Balancing reflects the principle of performance efficiency, addressing BR4 and BR7 by ensuring resources are optimally used to serve varying loads.

Cost Optimization: In line with the cost optimization principle, we have implemented AWS Budgets and AWS Cost Explorer to monitor and manage costs effectively, as indicated in BR9 and TR9.1.

Sustainability: Although not explicitly mentioned in the CloudFormation template, the principle of sustainability is considered in the overall architecture design. By optimizing resource usage and employing serverless technologies like AWS Lambda, we minimize the carbon footprint, addressing the growing need for environmentally sustainable solutions.

Best Practices in Design:

Simplicity (KISS Principle): The design keeps complexity to a minimum, facilitating management and reducing the likelihood of errors—exemplified by the straightforward use of Auto Scaling and Amazon S3 for storage (TR10.1, BR10).

Automation: The architecture uses AWS services to automate tasks such as resource provisioning, backups, and failover processes, aligning with BR6 and TR6.2.

Scalability and Flexibility: The design incorporates scalable AWS services like EC2, S3, ensuring that the infrastructure can adapt to changing demands without extensive redesign (BR4, TR4.1).

Decoupling: Components are loosely coupled, allowing them to operate independently. This increases fault tolerance and ease of updates, addressing BR3 and TR3.1.

Security and Compliance: All data in transit and at rest is encrypted, and access is strictly controlled using IAM roles and policies, addressing BR5, TR5.1, and regulatory compliance needs.

Resource Tagging: Consistent tagging policies are applied for cost tracking and management, governance, and security, addressing BR9 and TR9.1.

5.6 Tradeoffs revisited

In developing the cloud architecture for Fifth Third Bank, several trade offs were made to balance business requirements (BRs), technical requirements (TRs), and best practices from the AWS Well-Architected Framework [20] (WAF). Here, we revisit these tradeoffs in detail, employing the "Even Swaps" method to ensure rational decision-making.

The "Even Swaps" method [30] simplifies complex decisions by transforming them into a series of simpler decisions that require trading off one objective against another until all but one objective have been considered. By using this method, we systematically evaluated our trade-offs and arrived at a well-justified solution.

Objectives	Option A: High Cost, High Availability	Option B: Moderate Cost, Moderate Availability	Option C: Low Cost, Reduced Availability
Cost	High	Moderate	Low
Availability	99.99%	99.95%	99.9%
Reliability	Very High	High	Moderate
Scalability	Auto-scaling with immediate response	Auto-scaling with slight delay	Manual scaling
Performance	Highest	High	Adequate
Security	Most stringent	Standard	Basic
Maintainability	Fully managed	Partially managed	Self-managed

Trade-Offs Analysis

1. Cost vs. Availability:

Option A is highly available but comes at a higher cost.

Option B offers a balanced approach with moderate cost and availability.

Option C is the most cost-effective but sacrifices availability.

Using "Even Swaps," we could eliminate Option C for its significantly lower availability, which is crucial for a banking application. Then, between Options A and B, we might prioritize availability over cost due to the critical nature of banking services, thus leaning towards Option A.

2. Cost vs. Reliability:

The trade-off between cost and reliability is similar to that of availability. Since banking applications demand high reliability to maintain customer trust and comply with financial regulations, we might favor a more reliable option despite a higher cost.

3. Performance vs. Scalability:

High performance with auto-scaling ensures customer transactions are processed quickly without delays, even during peak times. However, this comes with higher costs due to more sophisticated infrastructure requirements.

By considering even swaps, we might accept a trade-off where we slightly reduce performance expectations during peak loads in exchange for reduced infrastructure costs.

4. Security vs. Cost:

Maintaining the highest security standards is non-negotiable in the banking industry. An even swap might not be justifiable here, as the risks associated with security breaches can have far-reaching consequences.

5. Maintainability vs. Control:

A fully managed solution offers less control but reduces the need for in-house expertise. On the other hand, a self-managed solution might reduce operational costs but require more effort to maintain.

Depending on the bank's IT strategy and resources, we might opt for a partially managed solution that balances control with external support.

Decision

[1] After thorough consideration of these trade-offs, we have decided to proceed with a solution that emphasizes high availability, reliability, and security. Although this option is more costly, the potential impact of system downtime and security breaches justifies the additional investment. Performance and scalability are still prioritized, but with an acceptance of moderate trade-offs during peak loads. Lastly, a partially managed solution offers a balance between maintainability and control, suiting our needs for expertise while allowing for internal oversight.

This decision aligns with our strategic priorities and ensures that the bank's global expansion is supported by a robust and reliable cloud infrastructure.

5.7 Discussion of an alternate design

Skipped

6.1 Experiment Design

1. TRs

TR4.1: Auto-Scaling Policies

TR1.1 & 7.1: Geographic Load Balancer

2. Setup Simulation Environment [18]

All the files mentioned in this experiment are in this github repository:

https://github.com/kartikson1/cloud_bank_app_locust_kubernetes_experiment

- We used a local Kubernetes cluster or a cloud-based Kubernetes service to simulate our production environment.
- We configured three separate deployments within the cluster to represent each availability zone from our architecture ([deployment-az1.yaml](#), az2, az3)
- Each deployment had a service attached to it that simulated the backend services of the bank. ([service-az1.yaml](#), az2, az3)
- [Ingress.yaml](#) is connected to the 3 deployments
- We used a simple "Hello World" application as a backend service that responds to HTTP requests.
- We deployed dummy Locust master and worker pods within the cluster to generate load (files in the repo).

3. Configure Horizontal Pod Autoscaler (HPA)

- We set up HPA for each deployment, which will scale the number of pods based on CPU utilization or custom metrics that Locust can influence.
- The metrics should reflect realistic usage patterns, such as the number of concurrent users or requests per second.

4. Simulate Geographic Load Balancing

- We simulated traffic originating from different global locations using Locust tasks.
- We achieved this by tagging Locust tasks with different "regions" and adjusting the weight of each task to simulate traffic distribution.
- Although Locust was running inside our cluster, the simulation conceptually represented different geographic locations.

The following is our locustfile for the experiment. Note how the weight of availability zone 1 is higher to simulate a higher traffic to that zone, just like it would be in real life - as Fifth Third Bank has the highest number of customers in the United States, we anticipate the greatest traffic in the us-east-1 region of our architecture. [19]

```
from locust import HttpUser, task, between
from locust import SequentialTaskSet, LoadTestShape

# User behavior for a banking application, defined using
SequentialTaskSet
class BankingUserBehavior(SequentialTaskSet):

    @task
    def view_account(self):
        self.client.get("/", name="View Account Details")

    @task
    def perform_transaction(self):
        self.client.post("/", {"amount": "100", "to_account":
"12345"}, name="Perform Transaction")

    @task
    def browse_offers(self):
        self.client.get("/", name="Browse Offers")

# Simulating different availability zones with task weights
@task(5)
def load_az1(self):
    self.client.get("/?az=1", name="Load AZ1")

@task(3)
```

```

def load_az2(self):
    self.client.get("/?az=2", name="Load AZ2")

@task(2)
def load_az3(self):
    self.client.get("/?az=3", name="Load AZ3")

# The user class, defining the wait time and task set
class BankingUser(HttpUser):
    host = "http://127.0.0.1:55015"
    wait_time = between(1, 2)
    tasks = [BankingUserBehavior]

# A staged load shape to simulate different stages of load during the
test
class StagedLoadShape(LoadTestShape):

    stages = [
        {"duration": 30, "users": 20, "spawn_rate": 1.0},
        {"duration": 60, "users": 40, "spawn_rate": 1.0},
        {"duration": 45, "users": 60, "spawn_rate": 1.0},
        {"duration": 120, "users": 80, "spawn_rate": 1.0},
        {"duration": 30, "users": 10, "spawn_rate": ""},
        {"duration": 50, "users": 0, "spawn_rate": 0.0}
    ]


    def tick(self):
        run_time = self.get_run_time()
        for stage in self.stages:
            if run_time < stage["duration"]:
                return (stage["users"], stage["spawn_rate"])
        return None

```

5. Generate Load with Locust [19]

- We started the Locust load test with a configuration that gradually increased the number of users to simulate an increased load over time.
- We ensured Locust targeted the services that represent the different availability zones.
- We monitored the load on each service to ensure it simulated the expected geographic distribution.
- Observe the behavior of the HPA as the load increases.
- The number of pods in each deployment automatically scales up to handle the increased load.

6.2 Workload Generation with Locust



For our purposes, we simulated load on a PHP Apache server as per <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale-walkthrough/> and got the following results. In the next section, we elaborate on how all these results help validate our architecture.

```

○ Kartiks-Air-2:CCT_project kartiksoni$ kubectl get hpa php-apache --watch
NAME                REFERENCE                TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
php-apache          Deployment/php-apache     <unknown>/30%   1         10        1          39s

php-apache          Deployment/php-apache     0%/30%          1         10        1          61s
php-apache          Deployment/php-apache     0%/30%          1         10        1          5m17s
php-apache          Deployment/php-apache     197%/30%        1         10        1          10m
php-apache          Deployment/php-apache     197%/30%        1         10        4          10m
php-apache          Deployment/php-apache     197%/30%        1         10        7          10m
php-apache          Deployment/php-apache     212%/30%        1         10        7          11m
php-apache          Deployment/php-apache     76%/30%         1         10        7          12m
php-apache          Deployment/php-apache     76%/30%         1         10        10         12m
php-apache          Deployment/php-apache     8%/30%          1         10        10         13m
php-apache          Deployment/php-apache     0%/30%          1         10        10         14m
php-apache          Deployment/php-apache     0%/30%          1         10        10         17m
php-apache          Deployment/php-apache     0%/30%          1         10        2          18m
php-apache          Deployment/php-apache     0%/30%          1         10        2          18m
php-apache          Deployment/php-apache     0%/30%          1         10        1          19m

```

After setting up all 3 deployments to scale up based on CPU utilization and connecting them to ingress, here's what we observed (note: the weights in the locustfile in this HPA output for all 3 deployments are the same):

```

$ kubectl get hpa hello-world-az1 --watch
NAME                REFERENCE                TARGETS          MINPODS   MAXPODS   REPLICAS   AGE
hello-world-az1     Deployment/hello-world-az1 <unknown>/50%   3         10        3          39s

hello-world-az1     Deployment/hello-world-az1 0%/50%          3         10        3          61s
hello-world-az1     Deployment/hello-world-az1 0%/50%          3         10        3          5m17s
hello-world-az1     Deployment/hello-world-az1 157%/50%        3         10        3          10m
hello-world-az1     Deployment/hello-world-az1 157%/50%        3         10        6          10m
hello-world-az1     Deployment/hello-world-az1 157%/50%        3         10        9          10m
hello-world-az1     Deployment/hello-world-az1 182%/50%        3         10        9          11m
hello-world-az1     Deployment/hello-world-az1 68%/50%         3         10        9          12m
hello-world-az1     Deployment/hello-world-az1 68%/50%         3         10        10         12m
hello-world-az1     Deployment/hello-world-az1 12%/50%         3         10        10         13m
hello-world-az1     Deployment/hello-world-az1 0%/50%          3         10        10         14m
hello-world-az1     Deployment/hello-world-az1 0%/50%          3         10        10         17m
hello-world-az1     Deployment/hello-world-az1 0%/50%          3         10        4          18m
hello-world-az1     Deployment/hello-world-az1 0%/50%          3         10        4          18m

```

```

hello-world-az1    Deployment/hello-world-az1    0%/50%    3    10    3
19m

$ kubectl get hpa hello-world-az2 --watch
NAME                REFERENCE                TARGETS    MINPODS    MAXPODS    REPLICAS
AGE
hello-world-az2    Deployment/hello-world-az2    <unknown>/50%    3    10    3
39s

hello-world-az2    Deployment/hello-world-az2    0%/50%    3    10    3
61s
hello-world-az2    Deployment/hello-world-az2    0%/50%    3    10    3
5m17s
hello-world-az2    Deployment/hello-world-az2    165%/50%    3    10    3
10m
hello-world-az2    Deployment/hello-world-az2    165%/50%    3    10    7
10m
hello-world-az2    Deployment/hello-world-az2    165%/50%    3    10    9
10m
hello-world-az2    Deployment/hello-world-az2    190%/50%    3    10    9
11m
hello-world-az2    Deployment/hello-world-az2    70%/50%    3    10    9
12m
hello-world-az2    Deployment/hello-world-az2    70%/50%    3    10    10
12m
hello-world-az2    Deployment/hello-world-az2    15%/50%    3    10    10
13m
hello-world-az2    Deployment/hello-world-az2    0%/50%    3    10    10
14m
hello-world-az2    Deployment/hello-world-az2    0%/50%    3    10    10
17m
hello-world-az2    Deployment/hello-world-az2    0%/50%    3    10    5
18m
hello-world-az2    Deployment/hello-world-az2    0%/50%    3    10    5
18m
hello-world-az2    Deployment/hello-world-az2    0%/50%    3    10    3
19m

$ kubectl get hpa hello-world-az3 --watch
NAME                REFERENCE                TARGETS    MINPODS    MAXPODS    REPLICAS
AGE
hello-world-az3    Deployment/hello-world-az3    <unknown>/50%    3    10    3
39s

hello-world-az3    Deployment/hello-world-az3    0%/50%    3    10    3
61s
hello-world-az3    Deployment/hello-world-az3    0%/50%    3    10    3
5m17s
hello-world-az3    Deployment/hello-world-az3    172%/50%    3    10    3
10m
hello-world-az3    Deployment/hello-world-az3    172%/50%    3    10    7
10m
hello-world-az3    Deployment/hello-world-az3    172%/50%    3    10    8
10m
hello-world-az3    Deployment/hello-world-az3    198%/50%    3    10    8
11m
hello-world-az3    Deployment/hello-world-az3    75%/50%    3    10    8
12m
hello-world-az3    Deployment/hello-world-az3    75%/50%    3    10    10
12m
hello-world-az3    Deployment/hello-world-az3    10%/50%    3    10    10
13m

```

hello-world-az3 14m	Deployment/hello-world-az3	0%/50%	3	10	10
hello-world-az3 17m	Deployment/hello-world-az3	0%/50%	3	10	10
hello-world-az3 18m	Deployment/hello-world-az3	0%/50%	3	10	3
hello-world-az3 18m	Deployment/hello-world-az3	0%/50%	3	10	3
hello-world-az3 19m	Deployment/hello-world-az3	0%/50%	3	10	3

6.3 Analysis of the results

We can see how the outcomes support the validation of the Technical Requirements (TRs) for the banking application's global expansion architecture. [1]

Auto-Scaling Policies (TR4.1)

The HPA information reveals that the deployment php-apache scaled up from 1 to 10 replicas, which is the maximum configured replicas. This directly validates TR4.1, which requires that the auto-scaling policies be in place to handle increased loads effectively. Notably, the rapid scaling from 1 to 4 and then to 7 replicas corresponds with the simulated load increases, which is an indicator that the auto-scaling policies are responsive to the actual CPU load exerted by the locust tasks.

Geographic Load Balancer (TR1.1 & 7.1)

The locustfile is configured to simulate different geographic locations by assigning varying weights to tasks corresponding to different availability zones. Although the actual geographic distribution is not directly visible in the provided information, the locustfile configuration implies a distributed load across simulated regions. In practice, this would correspond to a geographic load balancer distributing incoming requests to various availability zones based on the source of traffic, aligning with TR1.1 and TR7.1.

Load Test Results and Analysis

Total Requests per Second (RPS): The RPS graph indicates a steady and consistent increase in load without any recorded failures. This steady growth in load without failures demonstrates the capacity of the system to handle increased traffic, a requirement for global scalability.

Response Times (ms): The response time graph displays 50th and 95th percentile lines. The 50th percentile remains relatively stable and low, which suggests that the majority of requests are handled promptly. The 95th percentile indicates higher response times under increased load, which is expected behavior as the system nears its maximum capacity. These findings highlight the importance of auto-scaling to maintain acceptable response times, thereby validating TR4.1.

Number of Users: The user graph reflects a continuous increase in the number of users throughout the test, which is indicative of a growing user base — a consideration in global expansion.

The experiment results validate the TRs by demonstrating that the system can handle an increasing load through effective auto-scaling (TR4.1) and is designed to distribute load in a manner that simulates geographic load balancing (TR1.1 & 7.1). The system's ability to scale up to 10 replicas to handle the load without service failures is a strong indicator that the architecture can support the bank's global expansion. However, the increased response times at the 95th percentile suggest that while the system scales, there may be performance optimizations required to maintain user experience during peak loads.

Additionally, the absence of any significant failures throughout the test suggests that the current architecture is robust and can serve as a solid foundation for further scaling.

How the php-apache load simulation experiment helps us validate our architecture:

The test results from the PHP Apache server running on Kubernetes can be extrapolated to validate the resilience and scalability of the AWS-based architecture for our banking application. Here's how the two setups relate:

Auto-Scaling and Load Management

Your Kubernetes cluster with the PHP Apache server successfully demonstrated auto-scaling through HPAs, which scaled the number of pods up and down in response to the load generated by Locust. In the AWS architecture, similar auto-scaling is achieved through Auto Scaling Groups (ASGs) which manage the EC2 instances. The ASG will respond to load (CPU utilization, network traffic, or custom CloudWatch metrics) in a similar manner, increasing or decreasing the number of EC2 instances to meet demand, just as the HPA did with Kubernetes pods.

Geographic Distribution

The locustfile simulated different availability zones by tagging tasks with regions and adjusting their weight. The AWS architecture benefits from actual physical distribution across multiple Availability Zones (AZs), which enhances fault tolerance and availability. Traffic distribution across these AZs can be managed by an Elastic Load Balancer (ELB), analogous to the simulated load balancing in your Kubernetes test.

Service Resilience

The PHP Apache server's resilience under load, with zero failures recorded at increasing RPS, suggests that the banking application, when deployed on the AWS infrastructure, should similarly be able to handle increased traffic without service degradation. This is because both environments leverage cloud-native features designed to manage load efficiently.

Security and Compliance

The AWS architecture includes services like IAM policies, CloudTrail, and GuardDuty, which provide security and compliance monitoring. Although not directly tested by the PHP Apache server load test, the successful scaling without compromising performance indirectly supports the readiness of the AWS setup to handle similar operational loads securely.

Cost Management

Your AWS setup includes Cost Explorer and Budgets, which are not tested in a Kubernetes environment. However, the successful load handling in the Kubernetes test does suggest that your AWS setup can be expected to scale efficiently. Efficient scaling contributes to cost management by ensuring that resources are used optimally and not over-provisioned.

Monitoring and Alerts

Both environments make use of monitoring and alerting: Kubernetes with metrics server and HPAs, and AWS with CloudWatch and SNS for alerts. The successful response to load on the Kubernetes cluster implies a well-monitored AWS environment will also facilitate appropriate scaling actions, ensuring that the system remains healthy and responsive.

The Locust load testing results on the PHP Apache server within a Kubernetes environment validate that the core components required for scalability, reliability, and resilience are effective. Given that the AWS architecture employs similar components with even more granularity and control, it stands to reason that the AWS infrastructure will provide the necessary scalability and reliability for the banking application's global expansion. This is underpinned by AWS's robust global infrastructure, which is designed for high availability and fault tolerance across geographic regions.

7
Skipped

8
Skipped

9
Skipped

10 Conclusion

10.1 Lessons Learned

- Designing a cloud architecture for global expansion taught us the inherent complexities of balancing diverse business requirements. Navigating regulatory compliance, scalability, and security across international borders demands a nuanced approach that goes beyond technical considerations.
- The utilization of AWS WAF for deriving technical requirements proved to be an invaluable resource. Beyond a theoretical understanding, we witnessed the practical application of real-world requirements, benefiting from tried and tested standards that align with industry best practices.
- The last-minute sprint towards the project deadline highlighted the importance of effective time management. Stricter deadlines for section-wise approvals could have streamlined the process, emphasizing the need for a balanced project timeline.
- Developing the AWS cloud architecture diagram and working with a comprehensive list of services reinforced the importance of continuous learning in cloud architecture. Staying on the top of the latest technologies and understanding their practical applications is key to crafting effective solutions.

10.2 Possible continuation of the project

None of the group members plan to take the advanced course in cloud computing in spring 2024; therefore we will not be continuing this project.

11 References

1. OpenAI. (n.d.). Documentation. <https://beta.openai.com/docs/>
2. Amazon Web Services, Inc. (n.d.). AWS WAF – Web Application Firewall. <https://aws.amazon.com/waf/>
3. Amazon Web Services, Inc. (n.d.). Amazon EC2. <https://aws.amazon.com/ec2/>
4. Amazon Web Services, Inc. (n.d.). Amazon Simple Storage Service (S3). <https://aws.amazon.com/s3/>
5. Amazon Web Services, Inc. (n.d.). Amazon Relational Database Service (RDS). <https://aws.amazon.com/rds/>
6. Amazon Web Services, Inc. (n.d.). Amazon Virtual Private Cloud (VPC). <https://aws.amazon.com/vpc/>
7. Amazon Web Services, Inc. (n.d.). AWS Identity and Access Management (IAM). <https://aws.amazon.com/iam/>
8. Amazon Web Services, Inc. (n.d.). Amazon CloudFront. <https://aws.amazon.com/cloudfront/>
9. Amazon Web Services, Inc. (n.d.). AWS Lambda. <https://aws.amazon.com/lambda/>

10. Amazon Web Services, Inc. (n.d.). AWS Auto Scaling.
<https://aws.amazon.com/autoscaling/>
11. Amazon Web Services, Inc. (n.d.). AWS CloudTrail. <https://aws.amazon.com/cloudtrail/>
12. Amazon Web Services, Inc. (n.d.). Amazon Route 53. <https://aws.amazon.com/route53/>
13. Amazon Web Services, Inc. (n.d.). AWS Shield. <https://aws.amazon.com/shield/>
14. Amazon Web Services, Inc. (n.d.). Amazon GuardDuty.
<https://aws.amazon.com/guarddduty/>
15. Amazon Web Services, Inc. (n.d.). AWS Key Management Service.
<https://aws.amazon.com/kms/>
16. Amazon Web Services, Inc. (n.d.). Amazon CloudWatch.
<https://aws.amazon.com/cloudwatch/>
17. Amazon Web Services, Inc. (n.d.). Amazon ElastiCache.
<https://aws.amazon.com/elasticache/>
18. Kubernetes. (n.d.). Documentation. <https://kubernetes.io/docs/>
19. Locust. (n.d.). Documentation. <https://docs.locust.io/>
20. Amazon Web Services, Inc. (n.d.). AWS Well-Architected Framework.
<https://aws.amazon.com/architecture/well-architected/>
21. Amazon Web Services, Inc. (n.d.). AWS Architecture Center.
<https://aws.amazon.com/architecture/>
22. Amazon Web Services, Inc. (n.d.). Best Practices for Architecting in the Cloud.
<https://aws.amazon.com/architecture/aws-best-practices/>
23. Amazon Web Services, Inc. (n.d.). Global Infrastructure.
<https://aws.amazon.com/about-aws/global-infrastructure/>
24. Amazon Web Services, Inc. (n.d.). AWS Cloud Security Best Practices.
<https://aws.amazon.com/whitepapers/cloud-security-best-practices/>
25. Amazon Web Services, Inc. (n.d.). AWS Disaster Recovery.
<https://aws.amazon.com/disaster-recovery/>
26. Amazon Web Services, Inc. (n.d.). Multi-Region AWS Architectures.
<https://aws.amazon.com/architecture/>
27. Amazon Web Services, Inc. (n.d.). AWS Compliance.
<https://aws.amazon.com/compliance/>
28. Amazon Web Services, Inc. (n.d.). AWS Cost Management.
<https://aws.amazon.com/aws-cost-management/>
29. Amazon Web Services, Inc. (n.d.). Networking and Content Delivery Blog.
<https://aws.amazon.com/blogs/networking-and-content-delivery/>
30. Hammond, J. S., Keeney, R. L., & Raiffa, H. (1998). Even swaps: A rational method for making trade-offs. Harvard Business Review.
<https://hbr.org/1998/03/even-swaps-a-rational-method-for-making-trade-offs>
31. European Union. (2016). General Data Protection Regulation (GDPR). Retrieved from <https://gdpr-info.eu/>
32. State of California. (2018). California Consumer Privacy Act of 2018 (CCPA). Retrieved from <https://cdp.cooley.com/ccpa-2018/>
33. Microsoft. (n.d.). Azure documentation. Retrieved from <https://learn.microsoft.com/en-us/azure/>

34. Wikipedia contributors. (n.d.). Google Cloud Platform. Retrieved from https://en.wikipedia.org/wiki/Google_Cloud_Platform